




Parallel Computing with MATLAB

Shiran Golan

Application Engineer

shirang@systematics.co.il





Outline

- Parallel Computing Overview
- Parallel Computing Paradigm
 - Multicore Desktops
 - Cluster Hardware
- Programming Parallel Applications
- Using parfor Loops
- Parallel Computing Beyond Parfor



Parallel Computing and Acceleration

Approach	Options	Learn more on MATLAB Documentation
Apply best coding practices	<ul style="list-style-type: none">• Preallocation• Vectorization	Performance and Memory
Use explicit parallelism	<ul style="list-style-type: none">• CPU• GPU• Cluster	MATLAB Parallel Computing
Integrate with other languages	<ul style="list-style-type: none">• C• C++• Fortran	MATLAB API for Other Languages



Why Parallel Computing?

- **Why parallel computing?**
 - Need faster insight into your data
 - Computing infrastructure is broadly available (multicore desktops, clusters, gpus)
- **Why parallel computing with MATLAB?**
 - Accelerate workflows with minimal code changes
 - Focus on engineering and research, not computation



Utilizing Additional Processing Power

- **Built-in multithreading**


- Automatically enabled in MATLAB since R2008a
- Functions such as : fft, eig, svd, sort, etc.
- Leverage multicore CPUs

- **Parallel Computing Tools**

- High-level constructs let you parallelize MATLAB applications
- For a variety of applications
- Leverage CPUs, GPUs, and scale to clusters and clouds



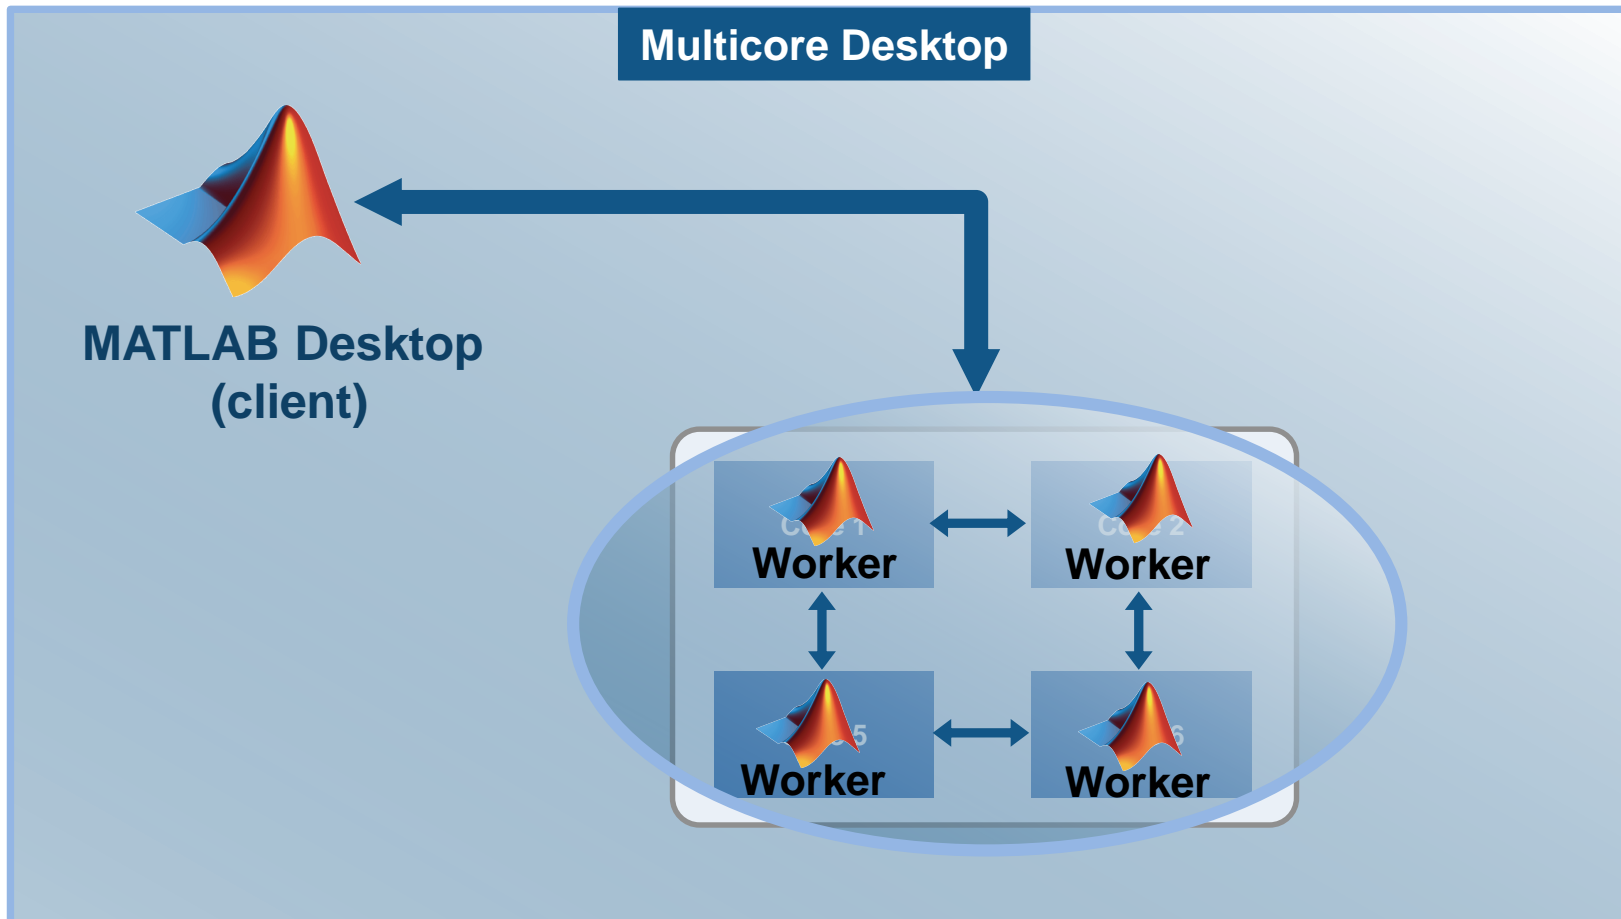
Outline

- Parallel Computing Overview
- Parallel Computing Paradigm
 -  - Multicore Desktops
 - Cluster Hardware
- Programming Parallel Applications
- Using parfor Loops
- Parallel Computing Beyond Parfor



Parallel Computing Paradigm

Multicore Desktops






Embarrassingly or Explicit Parallelism: Independent Tasks or Iterations

Parallel for loops

- Ideal problem for parallel computing
- No dependencies or communications between tasks
- Examples: parameter sweeps, Monte Carlo simulations







Parallel Computing Toolbox

Solve computationally and data-intensive problems using:

- multicore processors

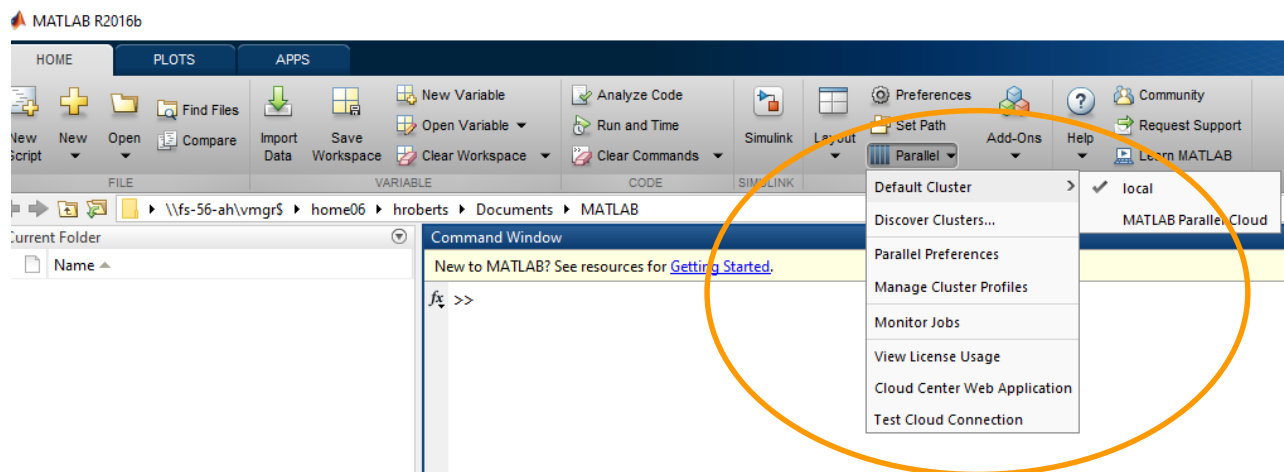


- computer clusters (requires MDCS)
- GPUs




Parallel Computing Toolbox

- Enables **explicit** use of multi-core processors
- Provides high-level tools and infrastructure for:
 - **creation** and **monitoring** of jobs
 - **allocating** hardware resources
- Provides a **bridge** to cluster hardware





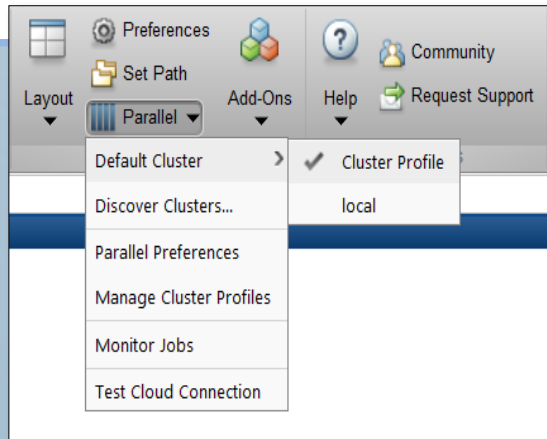
Outline

- Parallel Computing Overview
- Parallel Computing Paradigm
 - Multicore Desktops
 -  - Cluster Hardware
- Programming Parallel Applications + Example
- Using parfor Loops + Example
- Parallel Computing Beyond Parfor



Parallel Computing Paradigm

Cluster Hardware



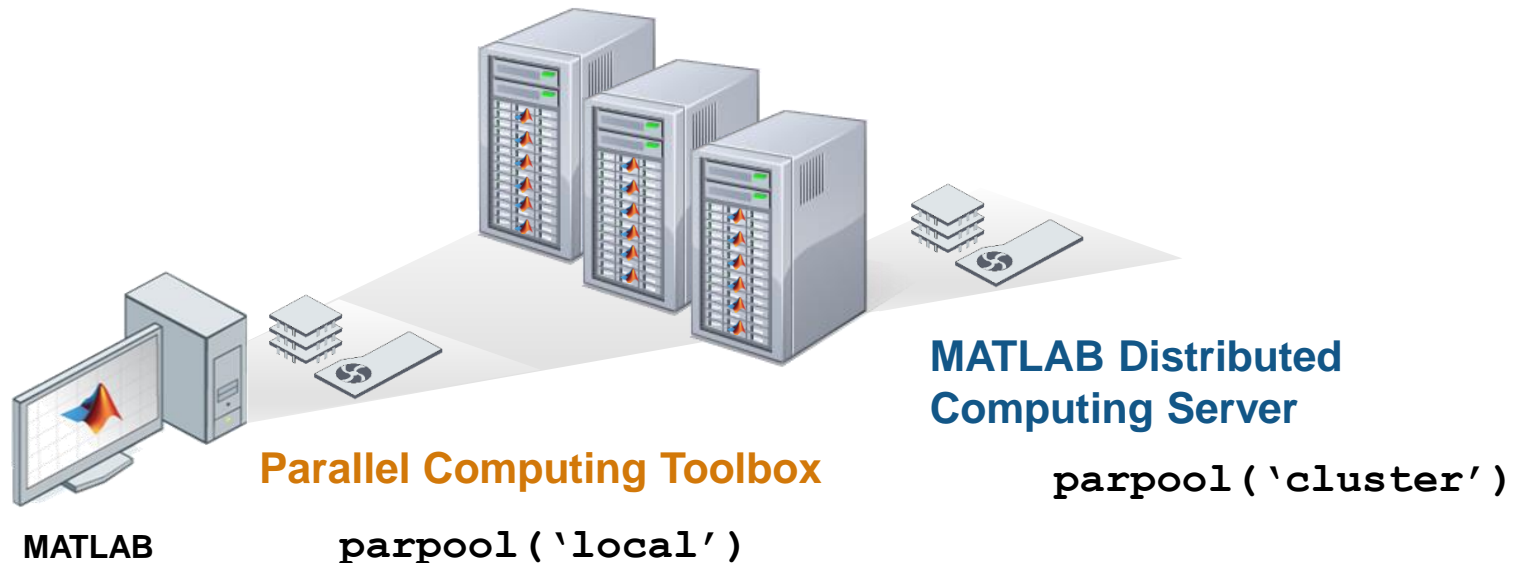

**MATLAB Desktop
(client)**

Cluster of Computers





Scale parfor



- Develop your application once
- Change run environment by changing the profile

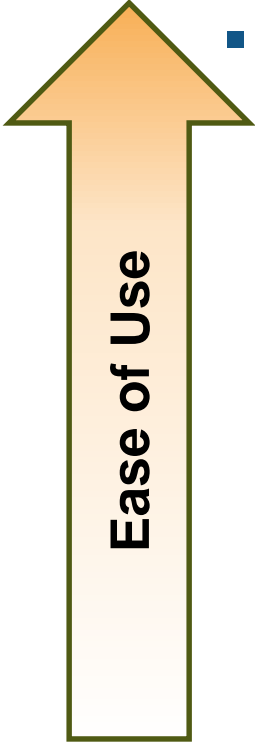


Outline

- Parallel Computing Overview
- Parallel Computing Paradigm
 - Multicore Desktops
 - Cluster Hardware
- Programming Parallel Applications
- Using parfor Loops
- Parallel Computing Beyond Parfor

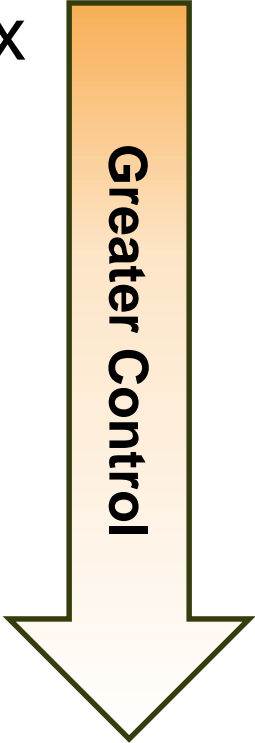


Programming Parallel Applications (1 of 3)



Ease of Use

- Enable with Parallel Computing Toolbox
 - Use same code with core MATLAB
 - Examples
 - Specific functions in many **toolboxes**
 - **parfor**
 - **batch**
 - **mapreduce**



Greater Control



Parallel-enabled Toolboxes

Enable parallel computing support by setting a flag or preference

Image Processing

Batch Image Processor, Block Processing, GPU-enabled functions




Original Image of Peppers



Recolored Image of Peppers


Statistics & Machine Learning

Resampling Methods, k-Means clustering, GPU-enabled functions




Neural Networks

Deep Learning, Neural Network training and simulation




Signal Processing and Communications

GPU-enabled FFT filtering, cross correlation, BER simulations




Computer Vision

Parallel-enabled functions in bag-of-words workflow



Optimization

Parallel estimation of gradients

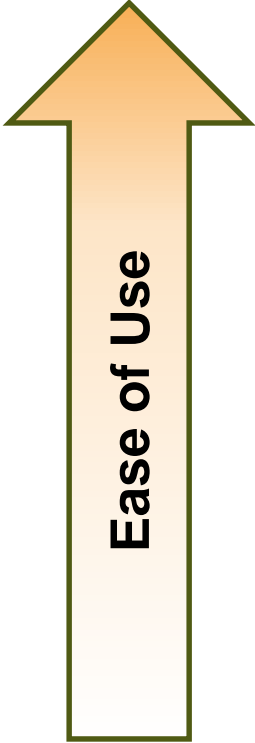


Other Parallel-enabled Toolboxes:

www.mathworks.com/products/parallel-computing/parallel-support.html



Programming Parallel Applications (2 of 3)



Ease of Use

- Enable by changing input data type
 - Use same algorithmic code with core MATLAB
 - Examples
 - Functions overloaded for **gpuArray**
 - Functions overloaded for **codistributed**

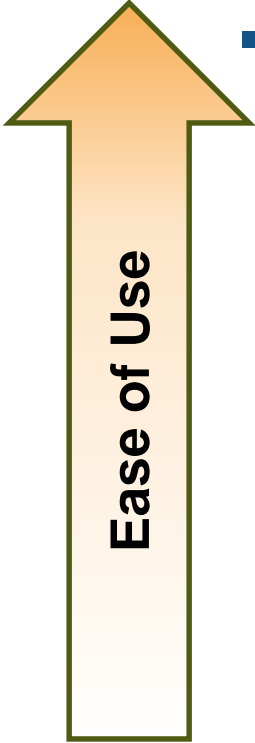


Greater Control

www.mathworks.com/help/distcomp/using-matlab-functions-on-codistributed-arrays.html
www.mathworks.com/help/distcomp/run-built-in-functions-on-a-gpu.html



Programming Parallel Applications (3 of 3)



Ease of Use

- Build code with Parallel Computing API
 - Applications require Parallel Computing Toolbox
 - Examples:
 - `spmd`, `gop`
 - Message passing
 - Job submission functions
 - CUDA kernels




Greater Control




Example: Optimizing Tower Placement

- Determine location of cell towers
- **Maximize** coverage
- **Minimize** overlap

Iteration: 1



Iteration: 22







Summary of Example

- Enabled built-in support for Parallel Computing Toolbox in Optimization Toolbox
- Used a **pool of MATLAB workers**
- Optimized in parallel using `fmincon`

Iteration: 1



Iteration: 22





Parallel Computing Support in Optimization Toolbox

- Functions:
 - `fmincon`

Finds a constrained minimum of a function of several variables
 - `fminimax`

Finds a minimax solution of a function of several variables
 - `fgoalattain`

Solves the multiobjective goal attainment optimization problem
- Functions can take finite differences in parallel in order to speed the estimation of gradients

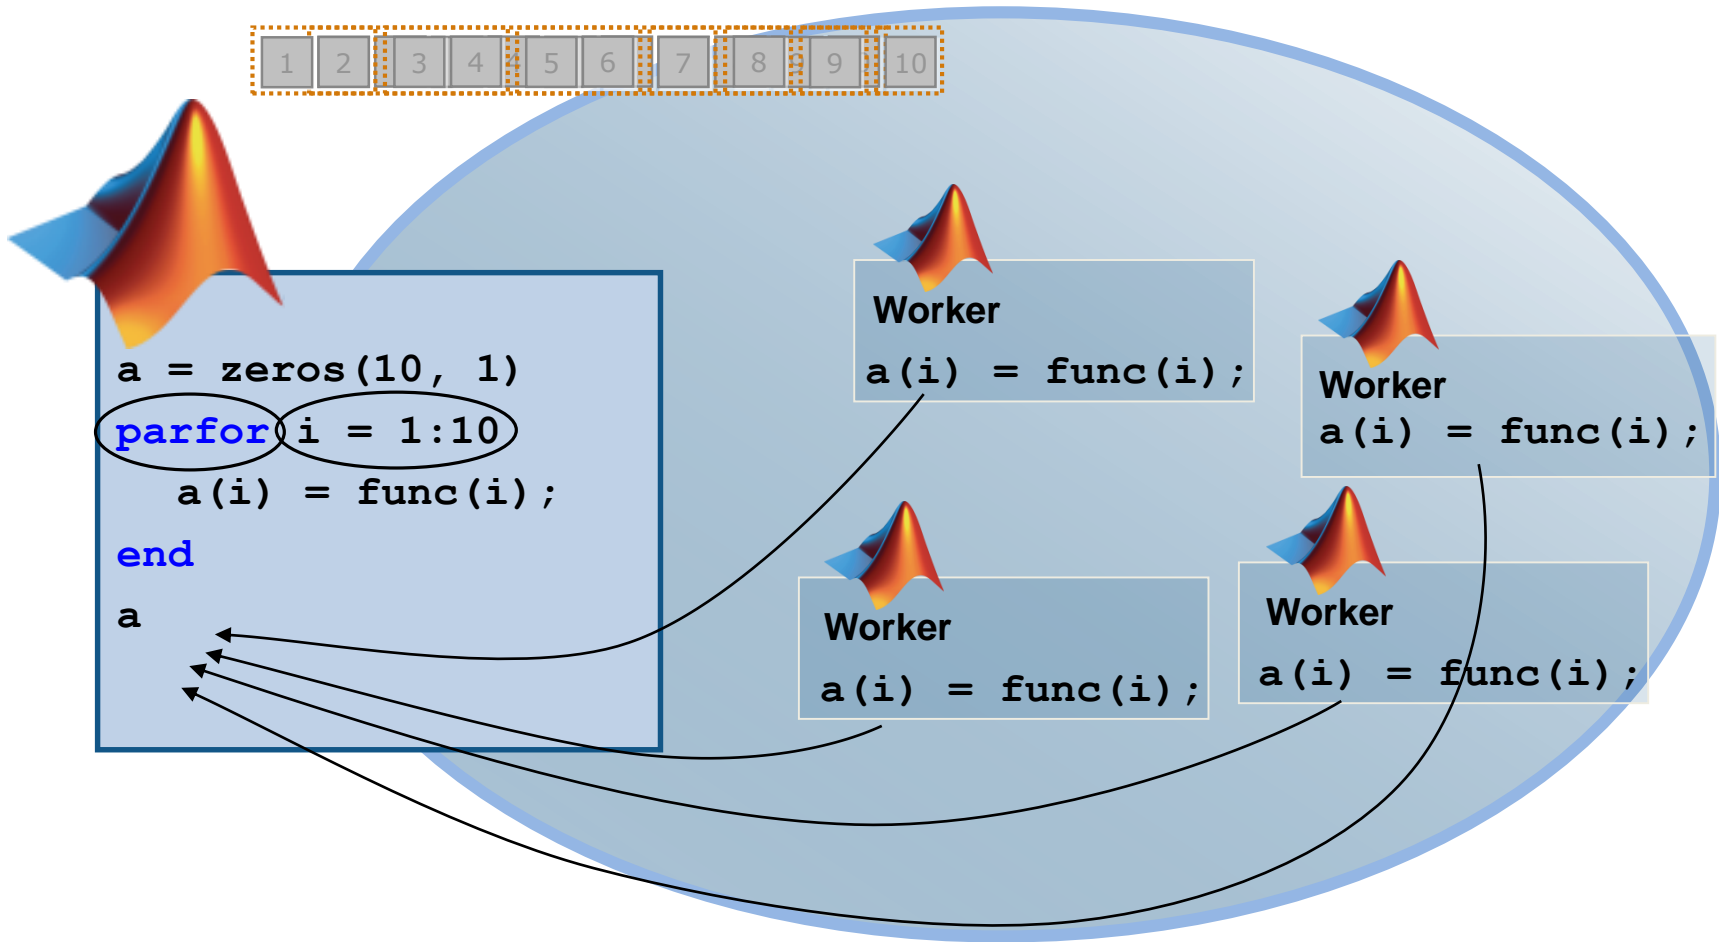


Outline

- Parallel Computing Overview
- Parallel Computing Paradigm
 - Multicore Desktops
 - Cluster Hardware
- Programming Parallel Applications
- Using parfor Loops
- Parallel Computing Beyond Parfor



The Mechanics of `parfor` Loops

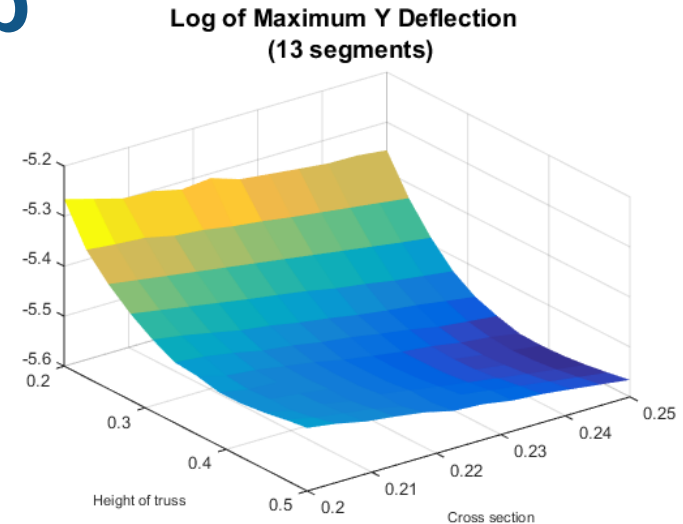


Pool of MATLAB Workers




Example: Parameter Sweep Parallel for-loops

- **Deflection** of customizable truss
 - Parameters investigated:
 - Height of truss
 - Cross sectional area of truss elements



- Convert `for` to `parfor`



- Use pool of MATLAB workers

$$M\ddot{x} + C\dot{x} + Kx = F$$

Parfeval: `N=3*p.NumWorkers;` -similar to what `parfor` does



Converting `for` to `parfor`

- Requirements for `parfor` loops:
 - Task independent
 - Order independent
- Constraints on the loop body:
 - Cannot “introduce” variables (e.g. `load`, etc.)
 - Cannot contain `break` or `return` statements
 - Cannot contain `evalc`, `eval`, `evalin`, `assignin`
 - Cannot contain a call to scripts directly
 - Cannot contain another `parfor` loop



Advice for Converting `for` to `parfor`

- Use **Code Analyzer** to diagnose `parfor` issues
- If your `for` loop cannot be converted to a `parfor`, consider **wrapping** a subset of the body to a function
- Read the section in the documentation on **classification** of variables
www.mathworks.com/help/distcomp/classification-of-variables-in-parfor-loops.html
- **Blog Post:** Using parfor Loops
blogs.mathworks.com/loren/2009/10/02/using-parfor-loops-getting-up-and-running/



Outline

- Parallel Computing Overview
- Parallel Computing Paradigm
 - Multicore Desktops
 - Cluster Hardware
- Programming Parallel Applications
- Using parfor Loops
- Parallel Computing Beyond Parfor



Parallel Computing with MATLAB – Beyond PARFOR

Well-known features

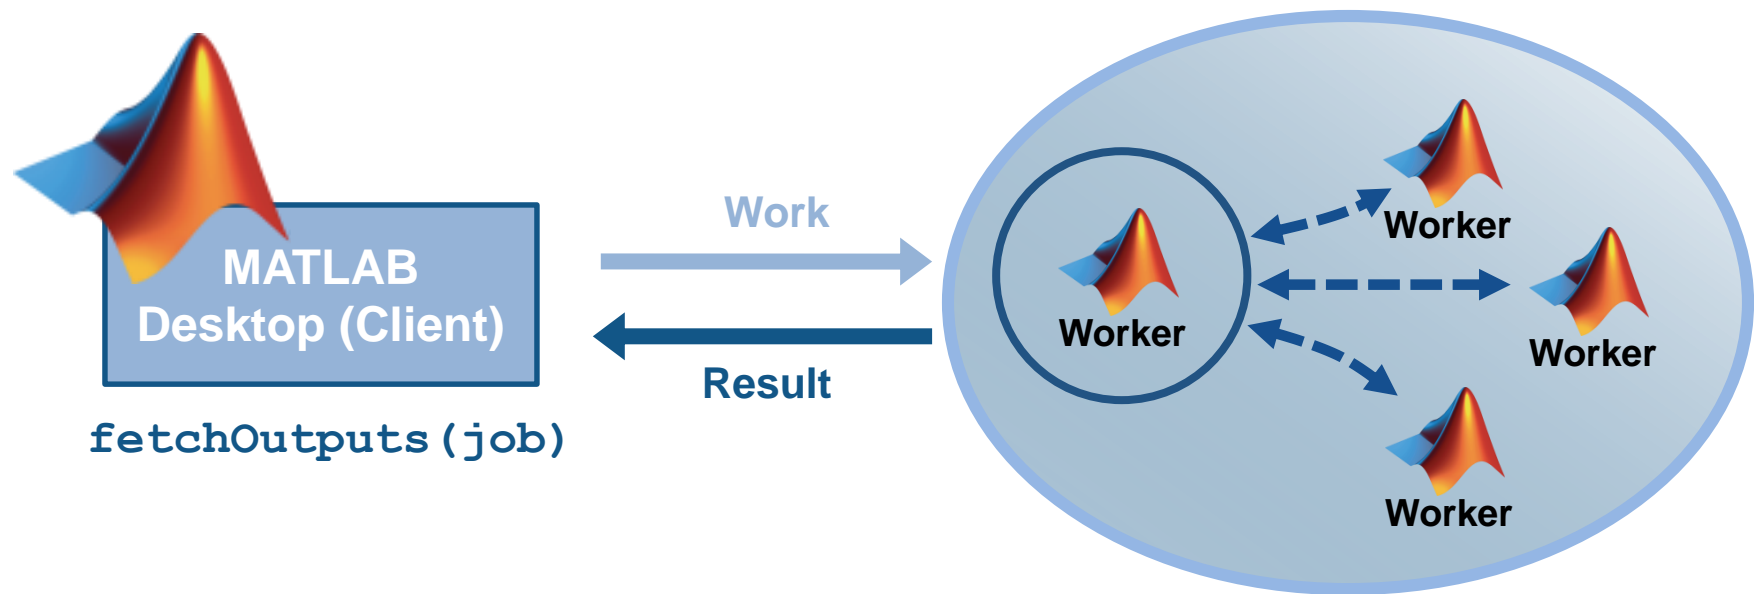
- parallel-enabled toolboxes
- `parfor`
- `gpuArray`

Advanced support

- batch submission, jobs and tasks
`batch`, `createJob`, `createTask`
- asynchronous execution on pool
`parfeval`
- parallel support for big data
`tall`, `mapreduce`
distributed arrays (“global arrays”) -
`distributed`, `codistributed`
- Single Program Multiple Data - `SPMD`
- message passing
`labSend`, `labReceive`



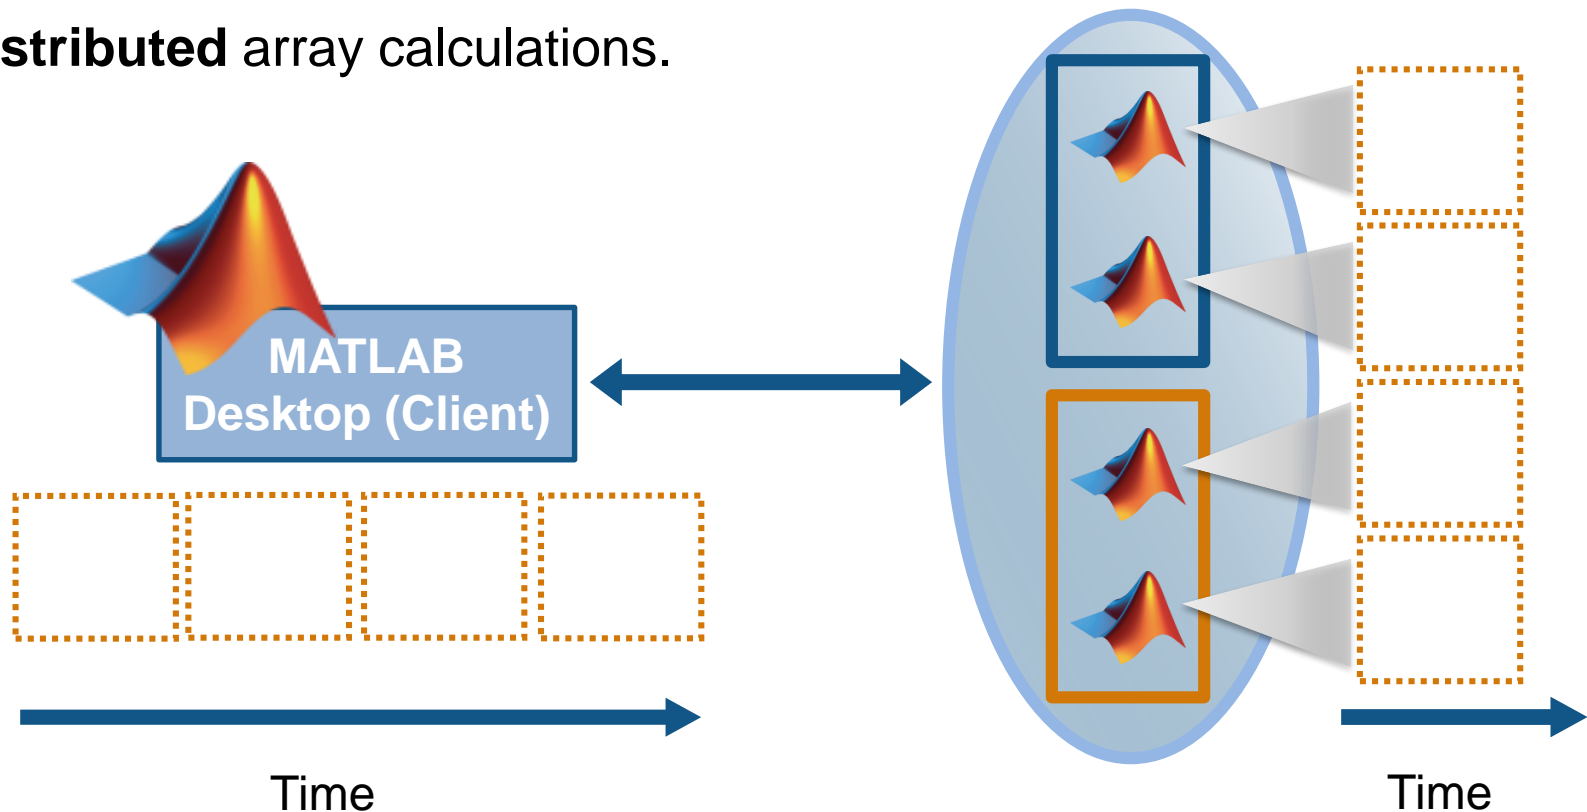
Offload and Scale Computations with batch






Execution with `spmd`

- A way to execute a block of code once on each worker in a pool.
- Can provide some time savings due to parallel execution, but it is most often used either to **set worker environments** or as a **building block for distributed** array calculations.





Execution with gop





Summary

- Parallel computing with MATLAB is **explicit parallelism** that uses **multiple MATLAB computation engines**
- Once you've written an application with **supported functions**, you can execute in parallel by **opening a parallel pool**
- Develop parallel applications on the desktop and then **scale** to clusters as needed



Speeding up MATLAB using the GPU



What is a GPU?

Graphical Processing Unit

NVIDIA GeForce GT 730M


# Cores	384, @719MHz
Memory	2 GB, DDR 3






History of GPUs

3D Gaming & CAD \longrightarrow Scientific Computing






CPUs vs. GPUs



CPU (Multiple Cores)



GPU (Hundreds of Cores)


System Memory

CPU	GPU
Sequential code, handles branching well	Simpler computational code
Large, low-latency memory	Small, high-latency memory
Several cores	Hundreds of cores
Can work on all data types	Strictly numerical data



Problems for Running on the GPU


A selection of problems from the CUDA Community Showcase:




Computational
Fluid Dynamics




Computational
Finance




Weather
Modeling



N-Body
Simulations



Molecular
Modeling



Digital Signal
Processing



The Language of GPUs

Kernel	Code running on the GPU
Device	The card containing the GPU and memory
Host	The CPU and system memory
CUDA	Compute Unified Device Architecture – Language designed for general purpose use of GPU



Three Levels of Support

- 1) Invoking built-in functions on arrays existing on the GPU
- 2) Translating simple MATLAB functions to run on the GPU
- 3) Invoking CUDA code from MATLAB




nVIDIA Solutions



Compute capability = 2.0 or higher (as of R2014b)



1. Invoking Built-In Functions






Built-In Functions That Support gpuArray

abs acos acosd acosh acot acotd acoth acsc acscd acsch accumarray all and angle any arrayfun asec asecd asech asin asind asinh assert atan atan2 atan2d atand atanh bandwidth besselj bessely beta betainc betaincinv betaln bicg bicgstab bitand bitcmp bitget bitor bitset bitshift bitxor blkdiag bsxfun cart2pol cart2sph cast cat cdf2rdf ceil chol circshift classUnderlying colon	companion complex cond conj conv conv2 convn corrcoef cos cosd cosh cot cotd coth cov cross csc cscd csch ctranspose cummax cummin cumprod cumsum deg2rad del2 det detrend diag diff discretize disp display dot double eig eps eq erf erfc erfcinv erfcx erfinv exp expint expm expm1 eye factorial false fft fft2 fftn fftshift filter filter2	find fix flip fliplr flipud floor fprintf full gamma gammaln gammalninv gammaln gather gc gmres gradient gt hankel histcounts horzcat hsv2rgb hypot idivide ifft ifft2 ifftn ifftshift imag ind2sub Inf inpolygon int16 int2str int32 int64 int8 interp1 interp2 interp3 interp intersect inv ipermute isaUnderlying isbanded iscolumn isdiag isempty isequal isequaln isfinite isfloat ishermitian isinf isinteger islogical	ismatrix ismember ismember_tol isnan isnumeric isreal isrow issorted issparse issyymmetric istrl istriu isvector kron ldivide le legendre length log log10 log1p log2 logical lsqr lt lu mat2str max median mean meshgrid min minus mldivide mod mode mpower mrdivide mtimes NaN ndgrid ndims ne nextpow2 nnz nonzeros norm normest not nthroot null num2str numel ones or orth	pagefun pcg perms permute pinv planerot plot (and related) plus pol2cart poly polyarea polyder polyfit polyint polyval polyvalm pow2 power prod psi qmr qr rad2deg rand randi randn randperm rank rdivide real reallog realpow realsqrt rectint rem repelem repmat reshape rgb2hsv roots rot90 round sec secd sech setdiff setxor shftdim sign sin sind single sinh size sort sortrows	spconvert sph2cart sprand sprandn sprandsym sprintf sqrt squeeze std sub2ind subsasgn subsindex subspace subsref sum superiorfloat svd swapbytes tan tand tanh times toeplitz trace transpose trapz tril triu true typecast uint16 uint32 uint64 uint8 uminus union unique unique_tol unwrap uplus vander var vertcat xor zeros
---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------


<https://www.mathworks.com/help/distcomp/run-built-in-functions-on-a-gpu.html>



Retrieving Data to the CPU







```
>> zCPU = gather(Z)
```





Summary: Invoking Built-In Functions





FFT on GPU

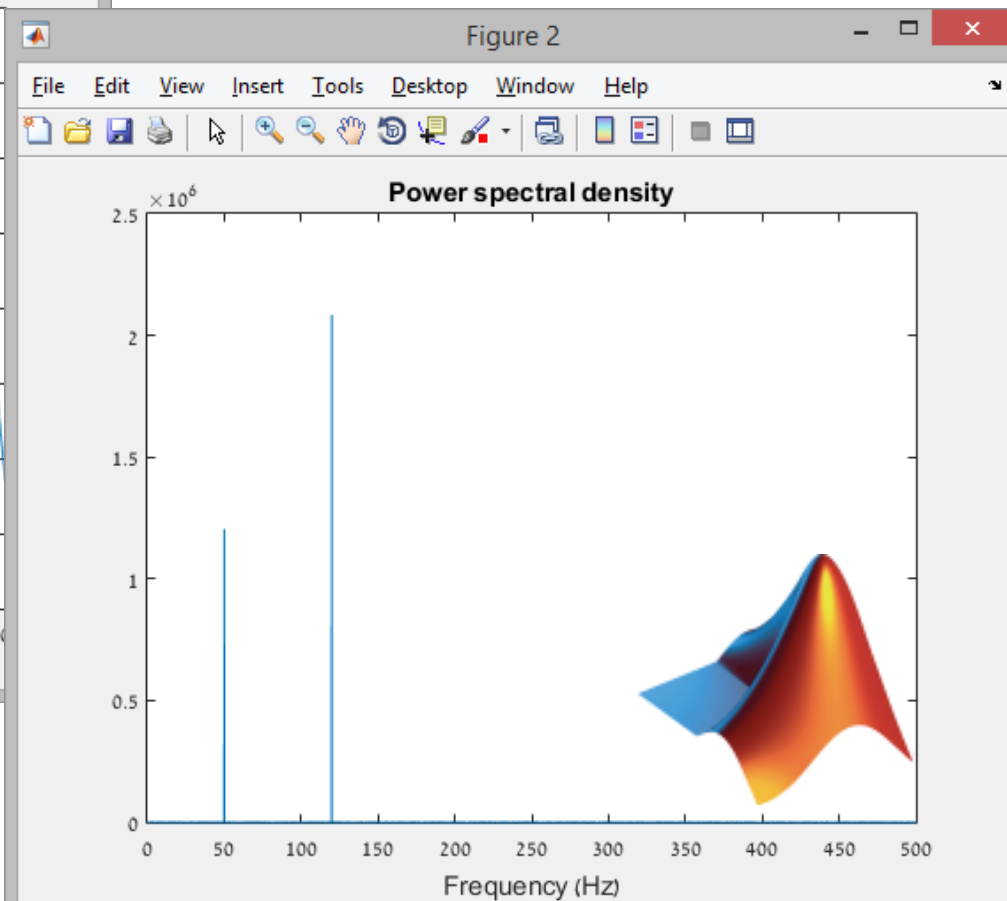
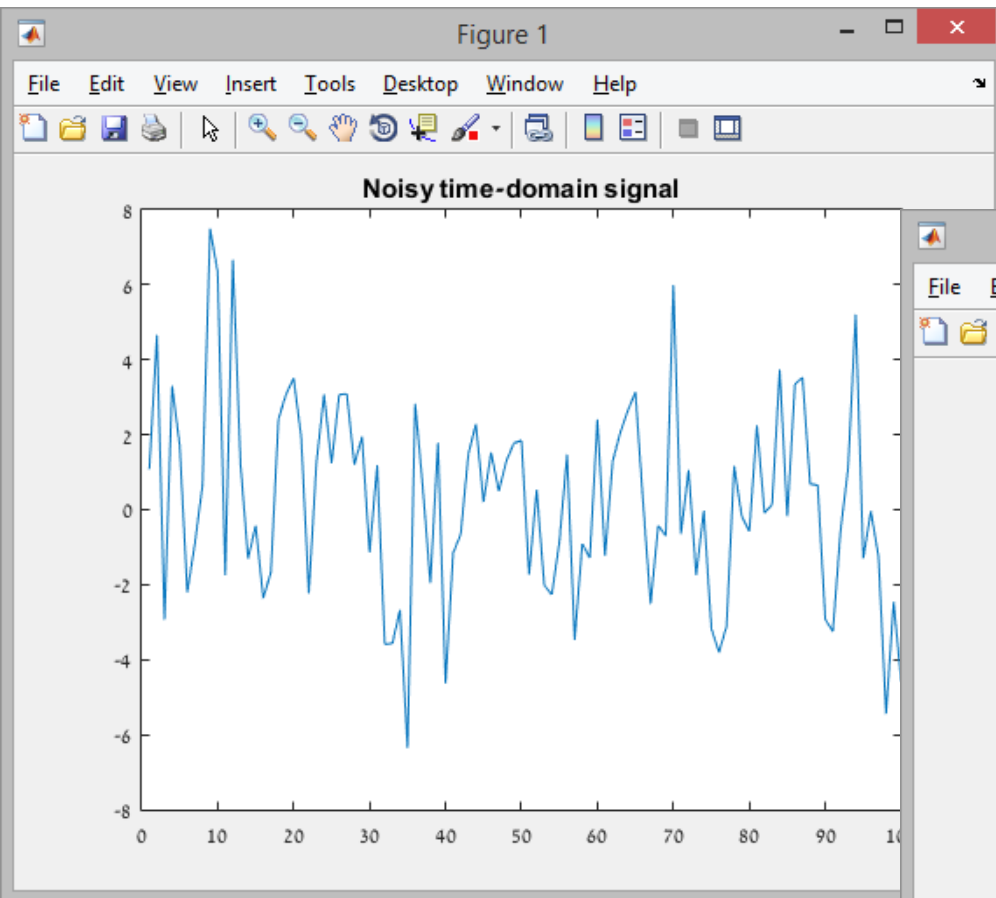




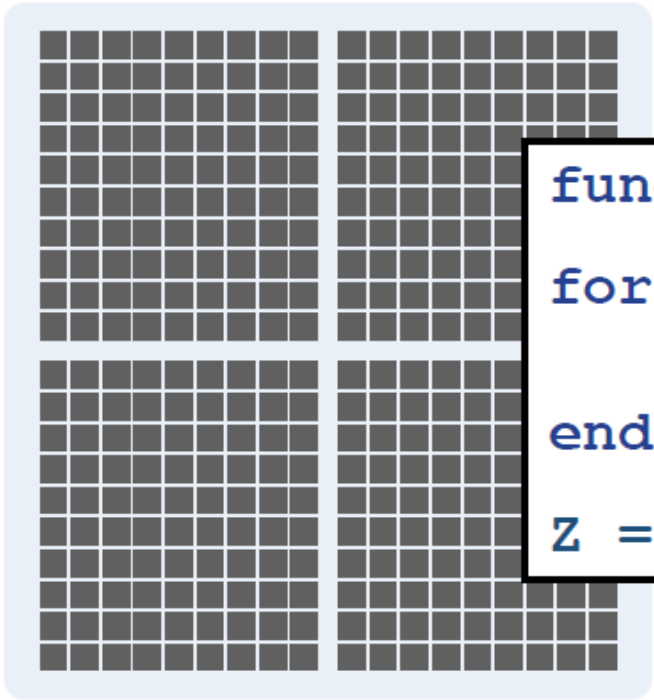
Image Processing & GPU

- GPU-enabled Image Processing Toolbox functions in [R2013a](#): *imrotate, imfilter, imdilate, imerode, imopen, imclose, imtophat, imbothat, imshow, padarray* and *bwlookup*
- 22 more functions in [R2013b](#), including: *edge, imresize, bwmorph* and *medfilt2*
- 9 more functions in [R2014a](#), including: *bwdist, imreconstruct, iradon, radon* and *imfill*
- *bwlabel* & *imregdemons* in [R2014b](#)
- 5 more in [R2015a](#), including *regionprops*
- 20 in [R2015b](#), including *bwareaopen, hough, imrotate, imresize* and *imcrop* (69 total)



2. Performing Calculations

```
>> Z = arrayfun(@juliaCalc,Z,-0.8+0.156i)
```




```
function Z = juliaCalc(Z,c)  
for k = 1:20  
    Z = Z.^2 + c;  
end  
  
Z = exp(-abs(Z));
```




Translating MATLAB Functions

- Accelerate scalar operations on large arrays
- Combine computations into a single kernel





3. Invoking a CUDA Kernel using feval



$$A_{circle} = \pi r^2$$

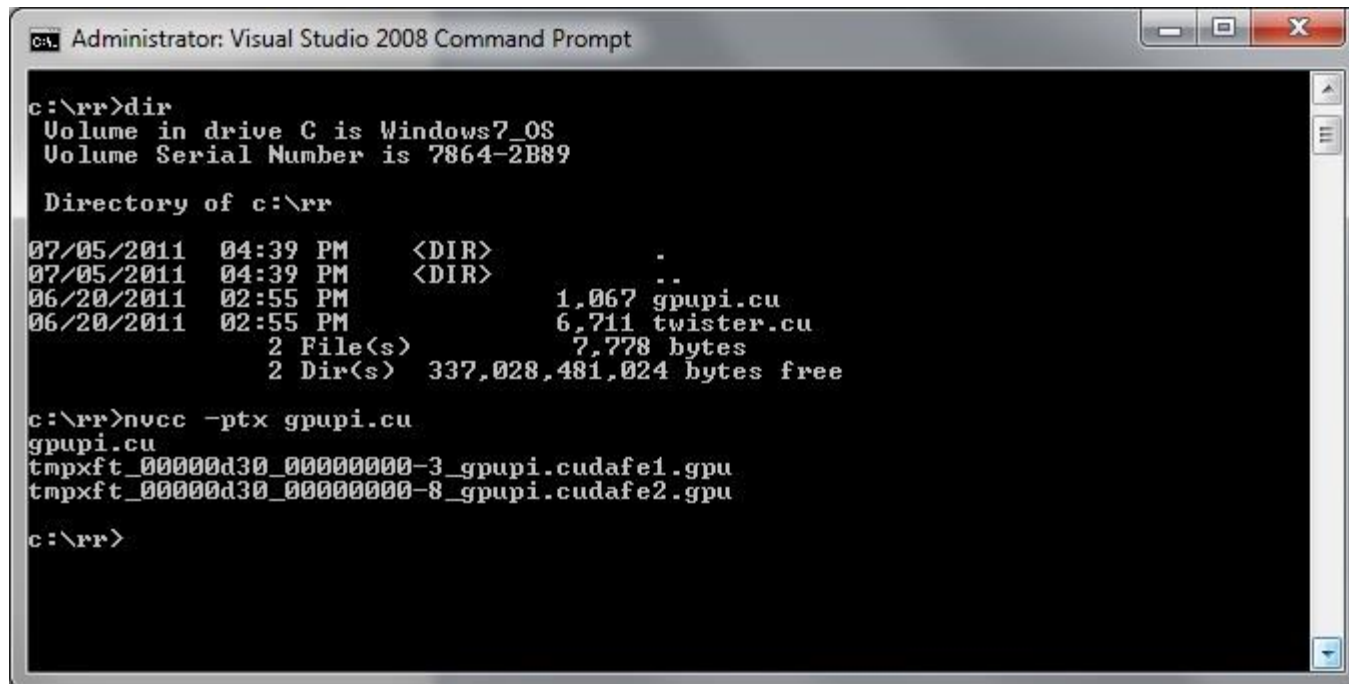
$$A_{square} = (2r)^2 = 4r^2$$

$$\frac{A_{circle}}{A_{square}} = \frac{\pi}{4}$$

$$P(x^2 + y^2 < r^2) = \frac{\pi}{4}$$



Generating Parallel Thread Execution Files



```
Administrator: Visual Studio 2008 Command Prompt

c:\rr>dir
Volume in drive C is Windows7_OS
Volume Serial Number is 7864-2B89

Directory of c:\rr

07/05/2011  04:39 PM    <DIR>          .
07/05/2011  04:39 PM    <DIR>          ..
06/20/2011  02:55 PM                1,067 gpupi.cu
06/20/2011  02:55 PM                6,711 twister.cu
               2 File(s)                7,778 bytes
               2 Dir(s)  337,028,481,024 bytes free

c:\rr>nvcc -ptx gpupi.cu
gpupi.cu
tmpxft_00000d30_00000000-3_gpupi.cudafe1.gpu
tmpxft_00000d30_00000000-8_gpupi.cudafe2.gpu

c:\rr>
```

Compile CUDA code to ptx using NVidia compiler




Summary: Invoking a CUDA Kernel

Pros

- Call existing CUDA code from MATLAB
- Highest level of functionality

Cons

- Requires knowledge of CUDA
- Requires knowledge of Parallel Programming





Summary of Features



Ease of Use

- 1) Invoking built-in functions on arrays existing on the GPU
- 2) Translating simple MATLAB functions to run on the GPU
- 3) Invoking CUDA code from MATLAB



Greater Functionality



2011-2016 Upgrades

- More **GPU-enabled** MATLAB functions & features, including capabilities from:
 - Statistics and Machine Learning Toolbox (>90)
 - Image Processing Toolbox (69)
 - Communications System Toolbox (11)
 - Signal Processing Toolbox (5)
 - Neural Network Toolbox
 - Phased Array System Toolbox
- MATLAB **Compiler** generated standalone executables and components now support applications that use the GPU



User Stories

NASA Langley Research Center Accelerates Acoustic Data Analysis with GPU Computing

The Challenge

Accelerate the analysis of sound recordings from wind tunnel tests of aircraft components

The Solution

Use MATLAB and Parallel Computing Toolbox to re-implement a legacy program for processing acoustic data, and cut processing time by running computationally intensive operations on a GPU

The Results




- Computations completed 40 times faster
- Algorithm GPU-enabled in 30 minutes
- Processing of test data accelerated




Wind tunnel test setup featuring the Hybrid Wing Body model (inverted), with 97-microphone phased array (top) and microphone tower (left).



More information

- Web:
www.mathworks.com/discovery/matlab-gpu.html
-  **MATLAB CENTRAL** (~900 results):
<http://www.mathworks.com/matlabcentral/>
- **MATLAB with Fun** Blog:
<http://matlabisrael.blogspot.com/>
- **Linked ** Group:
MATLAB & Simulink users in Israel 
- Seminars:
<http://www.systematics.co.il/products/mathworks/events/>
- Courses:
<http://www.systematics.co.il/courses/mathworks/>
- Webinars:
<http://www.mathworks.com/company/events/webinars/>
- Support & Sales: 03-7660111





Questions?

